

# The Alternating Direction Method of Multipliers

---

From the past, with vengeance

# Operator Splitting Quadratic Programming

We will tackle our QP using the OSQP solver by Oxford University

OSQP is a modern solver for Quadratic Programs in the form:

$$\arg \min_x \left\{ \frac{1}{2} x^T P x + q^T x \mid l \leq A x \leq u \right\}$$

The solver:

- is very fast, especially for problems with sparse matrices
- is available under a (very permissive) Apache 2.0 license
- has API for many programming languages

**The solver relies on the Alternating Direction Method of Multipliers (ADMM)**

- ...Plus a bunch of clever "tricks" to improve speed
- Here we will discuss only the basic ADMM, to provide **an intuition**

# The Alternating Direction Method of Multipliers

The ADMM solves numerical constrained optimization problems in the form:

$$\begin{aligned} & \operatorname{argmin} f(x) + g(z) \\ & \text{subject to: } Ax + Bz = c \end{aligned}$$

- Where  $f$  and  $g$  are assumed to be convex

**The methods relies on a so-called augmented Lagrangian**

This is a reformulation where the constraints are turned into penalty terms:

$$\mathcal{L}_\rho(x, z, \lambda) = f(x) + g(z) + \lambda^T (Ax + Bz - c) + \frac{1}{2} \rho \|Ax + Bz - c\|_2^2$$

- The algorithm idea is to optimize the augmented Lagrangian
- ...And to encourage constraint satisfaction via the penalty terms
- In practice, this is done by adjusting the multiplier vector  $\lambda$

# The Alternating Direction Method of Multipliers

## The ADMM operates as follows

We start from an initial assignment  $\mathbf{x}^0, \mathbf{z}^0, \boldsymbol{\lambda}^0$ , then we iterate:

$$\mathbf{x}^{k+1} = \operatorname{argmin}_{\mathbf{x}} \mathcal{L}_{\rho}(\mathbf{x}, \mathbf{z}^k, \boldsymbol{\lambda}^k)$$

$$\mathbf{z}^{k+1} = \operatorname{argmin}_{\mathbf{z}} \mathcal{L}_{\rho}(\mathbf{x}^{k+1}, \mathbf{z}, \boldsymbol{\lambda}^k)$$

$$\boldsymbol{\lambda}^{k+1} = \boldsymbol{\lambda}^k + \rho(\mathbf{A}\mathbf{x}^{k+1} + \mathbf{B}\mathbf{z}^{k+1} - \mathbf{c})$$

In other words:

- We keep everything fixed and we optimize over  $\mathbf{x}$  to obtain  $\mathbf{x}^{k+1}$
- We replace  $\mathbf{x}^k$  with  $\mathbf{x}^{k+1}$ , keep everything fixed and optimize over  $\mathbf{z}$
- Finally, we update the multiplier vector

**The switch between  $\mathbf{x}$  and  $\mathbf{z}$  optimization is the "alternating" part**

...While the use of the multipliers  $\boldsymbol{\lambda}$  explains the rest of the name

# Multiplier Update

Let's try to understand better the multiplier update

...Which consists in the rule:

$$\lambda^{k+1} = \lambda^k + \rho(Ax^{k+1} + Bz^{k+1} - c)$$

- The term  $Ax^{k+1} + Bz^{k+1} - c$  is just the current constraint violation
- ...In particular both its **amount** and **direction**

# Multiplier Update

Let's try to understand better the multiplier update

...Which consists in the rule:

$$\lambda^{k+1} = \lambda^k + \rho(Ax^{k+1} + Bz^{k+1} - c)$$

- The term  $Ax^{k+1} + Bz^{k+1} - c$  is just the current constraint violation
- ...In particular both its **amount** and **direction**

**If  $(Ax^{k+1} + Bz^{k+1})_i > c_i$  for some constraint  $i$ :**

- Then we **increase** the corresponding multiplier  $\lambda_i$
- So that the penalty term  $\lambda_i(Ax^{k+1} + Bz^{k+1} - c)_i$  grows
- This will push the next iteration to reduce the degree of violation

# Multiplier Update

Let's try to understand better the multiplier update

...Which consists in the rule:

$$\lambda^{k+1} = \lambda^k + \rho(Ax^{k+1} + Bz^{k+1} - c)$$

- The term  $Ax^{k+1} + Bz^{k+1} - c$  is just the current constraint violation
- ...In particular both its **amount** and **direction**

**If  $(Ax^{k+1} + Bz^{k+1})_i < c_i$  for some constraint  $i$ :**

- Then we **decrease** the corresponding multiplier  $\lambda_i$
- So that the penalty term  $\lambda_i(Ax^{k+1} + Bz^{k+1} - c)_i$  grows (again)
- This will push the next iteration to reduce the degree of violation (again)

# Multiplier Update

Let's try to understand better the multiplier update

...Which consists in the rule:

$$\lambda^{k+1} = \lambda^k + \rho(Ax^{k+1} + Bz^{k+1} - c)$$

- The term  $Ax^{k+1} + Bz^{k+1} - c$  is just the current constraint violation
- ...In particular both its **amount** and **direction**

**If  $(Ax^{k+1} + Bx^{k+1})_i = c_i$  for some constraint  $i$ :**

- Then we **keep** the corresponding multiplier  $\lambda_i$  **as it is**
- The constraint is not violated, so there is nothing to do



# Main Advantages of the Method

The method has two major advantages:

## 1) The $x$ and $z$ variables can be handled in isolation

- This results into simpler problems
- ...And in some cases enables massive parallelization

## 2) The ADMM converges under relatively mild conditions

- In the classical formulation,  $f$  and  $g$  need to be closed, proper, convex functions
  - They do not need to be differentiable
  - They can take the value  $+\infty$
  - We will see why that matters in the next slides
- The second condition is that  $\mathcal{L}_0(x, z, \lambda)$  should have a saddle point
  - This one is way trickier to check...

The full convergence proof can be found e.g. [here](#)

## **The ADMM and QP**

---

That was the whole point, right?

# QP Reformulation

Let's see these advantages at work on Quadratic Programs

We need to solve:

$$\operatorname{argmin}_x \left\{ \frac{1}{2} x^T P x + q^T x \mid l \leq A x \leq u \right\}$$

...Which we reformulate to:

$$\begin{aligned} & \operatorname{argmin} x^T P x + q^T x \\ & \text{subject to: } z = A x \\ & \quad l \leq z \leq u \end{aligned}$$

- We have introduced a new variables  $z$
- ...And posted the inequality constraints over that

# QP Reformulation

Then, we turn the inequality constraints into a function

$$\begin{aligned} & \operatorname{argmin} x^T P x + q^T x + \chi_{l \leq z \leq u}(z) \\ & \text{subject to: } z = A x \end{aligned}$$

Where:

- $\chi_{l \leq z \leq u}$  is the **characteristic function** of  $l \leq z \leq u$ 
  - It's value is  $+\infty$  when the constraint is violated and 0 elsewhere
  - In this case, it is non-differentiable, but closed, proper, and convex!
- $x^T P x + q^T x$  is our usual cost term
  - It is differentiable
  - ...And closed, proper, and convex if  $P$  is semi-definite positive

**We can now proceed to apply the ADMM!**

## The ADMM Steps for QP

**We need to start from a feasible  $x^0, z^0, \lambda^0$ :**

- That's easy, we get it by setting  $\lambda^0 = 0, z^0 = l$ , then solving  $Ax^0 = l$

**The  $x$  minimization step for  $\hat{z} = z^k$  becomes:**

$$\operatorname{argmin}_x x^T P x + q^T x + \chi_{l \leq z \leq u}(\hat{z}) + \lambda^T (\hat{z} - Ax) + \frac{1}{2} \rho \|\hat{z} - Ax\|_2^2$$

And then, since  $\hat{z}$  is fixed and feasible:

$$\operatorname{argmin} x^T P x + q^T x + \lambda^T (\hat{z} - Ax) + \frac{1}{2} \rho \|\hat{z} - Ax\|_2^2$$

This is a convex, differentiable, quadratic minimization problem

- It can be tackled via gradient descent
- ...Or by solving a linear system of equations

## The ADMM Steps for QP

The  $z$  minimization step for  $\hat{x} = x^{k+1}$  becomes

$$\operatorname{argmin}_z \hat{x}^T P \hat{x} + q^T \hat{x} + \chi_{l \leq z \leq u}(z) + \lambda^T (z - A \hat{x}) + \frac{1}{2} \rho \|z - A \hat{x}\|_2^2$$

Since  $\hat{x}$  is fixed, this can be reformulated as:

$$\begin{aligned} & \operatorname{argmin} \lambda^T z + \frac{1}{2} \rho \|z - A \hat{x}\|_2^2 \\ & \text{subject to: } l \leq z \leq u \end{aligned}$$

...And finally separated in to  $n$  problems (one per variable) in the form:

$$\operatorname{argmin}_{z_j} \left\{ \lambda_j z_j + \frac{1}{2} \rho (z_j - A_j \hat{x})^2 \mid l \leq z_j \leq u \right\}$$

## Some Considerations

### **We used the ADMM to break QP into a sequence of simpler problems**

The method can be used in other clever ways:

- Optimization with non-differentiable regularizers
- Parallel training, by splitting examples into multiple problems
- ...And using constraints to reach a consensus

### **The ADMM is best used for convex problems**

- Classical results are for convex problems only
- There are some (local) results for non-convex problems (e.g. [this one](#))
- ...But in practice it's less reliable

### **About the convergence pace**

- It's very fast in the first iterations, but much slower later
- You can get high-quality solutions early, but reaching the optimum takes long
- All in all, it's best to use the ADMM as an approximate method