

Symmetries

Sometimes metrics are not enough

Unexpected Discrepancy

Our current solution seems apparently perfect

```
In [14]: util.print_solution(tug, rflows, rpaths, sort='descending')
         sse = util.get_reconstruction_error(tug, rflows, rpaths, node_counts, arc_counts)
         print(f'RSSE: {np.sqrt(sse):.2f}')
```

```
8.17: 2,3 > 3,3
5.47: 0,2 > 1,2 > 2,2 > 3,2
3.74: 3,3
2.81: 0,1 > 1,1 > 2,0 > 3,0
2.09: 0,1 > 1,1 > 2,0 > 3,2
2.09: 1,0 > 2,0 > 3,0
1.24: 1,0 > 2,0 > 3,2
RSSE: 0.00
```

...And yet it **does not match** the ground truth!

```
In [7]: util.print_ground_truth(flows, paths, sort='descending')
```


```
8.17: 2,3 > 3,3
5.47: 0,2 > 1,2 > 2,2 > 3,2
4.89: 0,1 > 1,1 > 2,0 > 3,0
3.74: 3,3
3.32: 1,0 > 2,0 > 3,2
```

Unexpected Discrepancy

The discrepancy is unexpected, due to the 0 reconstruction error

Indeed, we can check that the reconstructed counts match the true ones:

```
In [15]: rnc, rac = util.get_counts(tug, rflows, rpaths)
util.plot_dict(rnc, figsize=figsize, label='reconstructed', data2=node_counts, label2='target',
```



Unexpected Discrepancy

What is going on?

Unexpected Discrepancy

What is going on?

We mentioned early on that **the available information is poor**

- There are many possible paths
- ...And many possible ways to explain the original counts!

How do we fix these **symmetries**?

Unexpected Discrepancy

What is going on?

We mentioned early on that **the available information is poor**

- There are many possible paths
- ...And many possible ways to explain the original counts!

How do we fix these **symmetries**?

- The only way is adding external information (e.g. a preference on paths)
- We can view this as a form of regularization

Occam's Razor

Intuitively, we could give priority to **the simplest explanation**



Image credit: [xkcd 2541](#)

A reasonable choice may be to use **a small number of paths**

How do we enforce this?

L1 Regularization and Path Number

We may think of using an L1 regularization

We would just need to add a linear term to the path formulation:

$$\arg \min_x \left\{ \frac{1}{2} x^T P x + q^T x + \alpha x \mid x \geq 0 \right\}$$

...Which would translate into a correction on the q vector:

$$\arg \min_x \left\{ \frac{1}{2} x^T P x + (q^T + \alpha) x \mid x \geq 0 \right\}$$

- This trick is implemented in the `solve_path_selection_full` function
- We just need to pass a value for the `alpha` argument

L1 Regularization and Path Number

Let's begin by trying $\alpha = 1$

```
In [16]: rflows2, rpaths2 = util.solve_path_selection_full(tug, node_counts, arc_counts, verbose=0, alpha=1)
print('FLOW: PATH')
util.print_solution(tug, rflows2, rpaths2, sort='descending')
sse = util.get_reconstruction_error(tug, rflows2, rpaths2, node_counts, arc_counts)
print(f'\nRSSE: {np.sqrt(sse):.2f}')
```

```
FLOW: PATH
8.10: 2,3 > 3,3
5.37: 0,2 > 1,2 > 2,2 > 3,2
2.58: 0,1 > 1,1 > 2,0 > 3,0
2.36: 3,3
1.98: 1,0 > 2,0 > 3,0
1.90: 0,1 > 1,1 > 2,0 > 3,2
1.17: 1,0 > 2,0 > 3,2
0.36: 0,1 > 1,1 > 2,0 > 3,3
0.06: 1,0 > 2,3 > 3,3
0.02: 0,1 > 1,0 > 2,0 > 3,0
0.02: 1,0 > 2,0 > 3,3
```

```
RSSE: 1.30
```

L1 Regularization and Path Number

Let's begin by trying $\alpha = 1$

```
In [16]: rflows2, rpaths2 = util.solve_path_selection_full(tug, node_counts, arc_counts, verbose=0, alpha=1)
print('FLOW: PATH')
util.print_solution(tug, rflows2, rpaths2, sort='descending')
sse = util.get_reconstruction_error(tug, rflows2, rpaths2, node_counts, arc_counts)
print(f'\nRSSE: {np.sqrt(sse):.2f}')
```

```
FLOW: PATH
8.10: 2,3 > 3,3
5.37: 0,2 > 1,2 > 2,2 > 3,2
2.58: 0,1 > 1,1 > 2,0 > 3,0
2.36: 3,3
1.98: 1,0 > 2,0 > 3,0
1.90: 0,1 > 1,1 > 2,0 > 3,2
1.17: 1,0 > 2,0 > 3,2
0.36: 0,1 > 1,1 > 2,0 > 3,3
0.06: 1,0 > 2,3 > 3,3
0.02: 0,1 > 1,0 > 2,0 > 3,0
0.02: 1,0 > 2,0 > 3,3
```

```
RSSE: 1.30
```

- The RSSE grows (as it could be expected)
- But we have **more** paths!

L1 Regularization and Path Number

What if we make α larger?

```
In [17]: rflows2, rpaths2 = util.solve_path_selection_full(tug, node_counts, arc_counts, verbose=0, alpha=0.01)
print('FLOW: PATH')
util.print_solution(tug, rflows2, rpaths2, sort='descending')
sse = util.get_reconstruction_error(tug, rflows2, rpaths2, node_counts, arc_counts)
print(f'\nRSSE: {np.sqrt(sse):.2f}')
```

FLOW: PATH

4.76: 2,3 > 3,3

4.27: 0,2 > 1,2 > 2,2 > 3,2

1.83: 0,1 > 1,1 > 2,0 > 3,0

1.42: 0,1 > 1,1 > 2,0 > 3,2

0.84: 0,1 > 1,1 > 2,0 > 3,3

0.82: 1,0 > 2,3 > 3,3

0.77: 1,0 > 2,0 > 3,0

0.29: 1,0 > 2,0 > 3,2

0.19: 0,1 > 1,0 > 2,3 > 3,3

0.15: 0,1 > 1,0 > 2,0 > 3,0

0.06: 0,1 > 1,0 > 2,0 > 3,2

0.04: 1,0 > 2,0 > 3,3

0.04: 0,1 > 1,0 > 2,0 > 3,3

0.02: 0,0 > 1,0 > 2,0 > 3,2

RSSE: 9.11

L1 Regularization and Path Number

What if we make α larger?

```
In [17]: rflows2, rpaths2 = util.solve_path_selection_full(tug, node_counts, arc_counts, verbose=0, alpha=0.01)
print('FLOW: PATH')
util.print_solution(tug, rflows2, rpaths2, sort='descending')
sse = util.get_reconstruction_error(tug, rflows2, rpaths2, node_counts, arc_counts)
print(f'\nRSSE: {np.sqrt(sse):.2f}')
```

FLOW: PATH

4.76: 2,3 > 3,3

4.27: 0,2 > 1,2 > 2,2 > 3,2

1.83: 0,1 > 1,1 > 2,0 > 3,0

1.42: 0,1 > 1,1 > 2,0 > 3,2

0.84: 0,1 > 1,1 > 2,0 > 3,3

0.82: 1,0 > 2,3 > 3,3

0.77: 1,0 > 2,0 > 3,0

0.29: 1,0 > 2,0 > 3,2

0.19: 0,1 > 1,0 > 2,3 > 3,3

0.15: 0,1 > 1,0 > 2,0 > 3,0

0.06: 0,1 > 1,0 > 2,0 > 3,2

0.04: 1,0 > 2,0 > 3,3

0.04: 0,1 > 1,0 > 2,0 > 3,3

0.02: 0,0 > 1,0 > 2,0 > 3,2

RSSE: 9.11

L1 Regularization and Path Number

Shouldn't L1 norm work as a sparsifier?

Not exactly: it simply results in a **fixed penalty rate** for raising a variable

- The solver will try to **balance** it with a larger reduction of the quadratic loss
- ...Which we can easily improve by including **more nodes** in each path

L1 Regularization and Path Number

Shouldn't L1 norm work as a sparsifier?

Not exactly: it simply results in a **fixed penalty rate** for raising a variable

- The solver will try to **balance** it with a larger reduction of the quadratic loss
- ...Which we can easily improve by including **more nodes** in each path

The truth is that when we use an L1 norm as sparsifier...

...We really wished our regularizer to be:

$$N_{paths} = \sum_{j=1}^n z_j \quad \text{with: } z_j = \begin{cases} 1 & \text{if } x_j > 0 \\ 0 & \text{otherwise} \end{cases}$$

- Which is inconvenient, since it is non-differentiable
- ...But what if we used an approach for non-differentiable optimization?

Path Consolidation Problem

Let's face an inconvenient truth

Path Consolidation Problem

For example, we could **focus on the paths in the current solution**:

- ...Minimize the number of used paths
- ...While preserving our reconstruction error

This is form of **symmetry breaking** (as a post-processing step)

By doing this, we obtain a **"path consolidation problem"** in the form:

$$\begin{aligned} & \arg \min_x \|z\|_1 \\ & \text{subject to: } Vx = v^* \\ & \quad Ex = e^* \\ & \quad x \leq Mz \\ & \quad x \geq 0 \\ & \quad z \in \{0, 1\}^n \end{aligned}$$

Path Consolidation Problem

Let's proceed to examine the formulation a bit better:

$$\begin{aligned} & \arg \min_x \|z\|_1 \\ & \text{subject to: } Vx = v^* \\ & \quad Ex = e^* \\ & \quad x \leq Mz \\ & \quad x \geq 0 \\ & \quad z \in \{0, 1\}^n \end{aligned}$$

- The terms V , E , and x are the same as before
- ...Except in this case we will consider a **a subset of the paths**
- v^* and e^* are the counts from the optimal path formulation solution
- We are requiring the (reconstructed) counts to be **exactly the same**

Path Consolidation Problem

Let's proceed to examine the formulation a bit better:

$$\begin{aligned} & \arg \min_x \|z\|_1 \\ & \text{subject to: } Vx = v^* \\ & \quad Ex = e^* \\ & \quad x \leq Mz \\ & \quad x \geq 0 \\ & \quad z \in \{0, 1\}^n \end{aligned}$$

- The z variables determine whether a path is used ($z_j = 1$) or not ($z_j = 0$)
- M is a constant large enough to make the constraint trivial if $z_j = 1$
- Constants such as these are often referred to as "big-Ms"
- Basically, $x \leq Mz$ is a linearization of the implication $x > 0 \Rightarrow z = 1$

Path Consolidation Problem

Let's proceed to examine the formulation a bit better:

$$\begin{aligned} & \arg \min_x \|z\|_1 \\ & \text{subject to: } Vx = v^* \\ & \quad Ex = e^* \\ & \quad x \leq Mz \\ & \quad x \geq 0 \\ & \quad z \in \{0, 1\}^n \end{aligned}$$

- All constraints are linear
- The cost function is linear
- Some variables are integer

This is a **Mixed Integer Linear Program (MILP)**